


14,672,234 members

Sign in [articles](#) [quick answers](#) [discussions](#) [features](#) [community](#) [help](#)[Articles » Platforms, Frameworks & Libraries » STL » Collections and Iterators](#)

# An In-Depth Study of the STL Deque Container

**Nitron**

11 Nov 2003

Rate this:  4.98 (158 votes)

This article presents an in-depth analysis of `std::deque` and offers guidance as to when to prefer using it as opposed to `std::vector`, by taking into consideration memory allocation and container performance.

## Introduction

This article presents an in-depth look at the STL **deque** container. This article will discuss the benefits of the **deque** and under what circumstances you would use it instead of the **vector**. After reading this article, the reader should be able to explain the fundamental differences between **vector** and **deque** with respect to container growth, performance and memory allocation. Since **deque** is so similar in usage and syntax to **vector**, it is recommended that the reader refers to this [article](#) [^] on **vector** for information on implementing and using this STL container.

## Deque Overview

The **deque**, like the **vector**, is also part of the Standard Template Library, the STL. The **deque**, or "double-ended queue", is very similar to the **vector** on the surface, and can act as a direct replacement in many implementations. Since it is assumed that the reader already knows how to effectively use the STL **vector** container, I have provided the table of **deque** member functions and operators below, solely for comparison and reference.

### Deque Member Functions<sup>1</sup>

Function	Description
<b>assign</b>	Erases elements from a <b>deque</b> and copies a new set of elements to the target <b>deque</b> .
<b>at</b>	Returns a reference to the element at a specified location in the <b>deque</b> .
<b>back</b>	Returns a reference to the last element of the <b>deque</b> .

<code>begin</code>	Returns an iterator addressing the first element in the <code>deque</code> .
<code>clear</code>	Erases all the elements of a <code>deque</code> .
<code>deque</code>	Constructs a <code>deque</code> of a specific size or with elements of a specific value or with a specific allocator or as a copy of all or part of some other <code>deque</code> .
<code>empty</code>	Tests if a <code>deque</code> is empty.
<code>end</code>	Returns an iterator that addresses the location succeeding the last element in a <code>deque</code> .
<code>erase</code>	Removes an element or a range of elements in a <code>deque</code> from specified positions.
<code>front</code>	Returns a reference to the first element in a <code>deque</code> .
<code>get_allocator</code>	Returns a copy of the allocator object used to construct the <code>deque</code> .
<code>insert</code>	Inserts an element or a number of elements or a range of elements into the <code>deque</code> at a specified position.
<code>max_size</code>	Returns the maximum length of the <code>deque</code> .
<code>pop_back</code>	Deletes the element at the end of the <code>deque</code> .
<code>pop_front</code>	Deletes the element at the beginning of the <code>deque</code> .
<code>push_back</code>	Adds an element to the end of the <code>deque</code> .
<code>push_front</code>	Adds an element to the beginning of the <code>deque</code> .
<code>rbegin</code>	Returns an iterator to the first element in a reversed <code>deque</code> .
<code>rend</code>	Returns an iterator that points just beyond the last element in a reversed <code>deque</code> .
<code>resize</code>	Specifies a new size for a <code>deque</code> .
<code>size</code>	Returns the number of elements in the <code>deque</code> .
<code>swap</code>	Exchanges the elements of two <code>deques</code> .

## Deque Operators<sup>1</sup>

Function	Description
<code>operator[]</code>	Returns a reference to the <code>deque</code> element at a specified position.

This striking similarity to `vector` then gives rise to the following question:

**Q:** If deque and vector offer such similar functionality, when is it preferable to use one over the other?

**A:** If you have to ask, use vector.

*Um, okay... How about a little explanation please?*

Why certainly, I'm glad you asked! I didn't pull that answer out of thin air, in fact, the answer comes from the C++ Standard<sup>2</sup> itself. Section 23.1.1 states the following:

*`vector` is the type of sequence that should be used by default. ... `deque` is the data structure of choice when most insertions and deletions take place at the beginning or at the end of the sequence.*

Interestingly enough, this article is practically devoted to fully understanding that statement.

## What's New

After perusing the table above and comparing it with `vector`, you will notice two new member functions.

1. `push_front()` - Adds elements to the front of a `deque`.
2. `pop_front()` - Removes elements from the front of a `deque`.

These are called with the same syntax as `push_back()` and `pop_back()`. Therein lies the first feature that could possibly warrant the use of `deque`, namely, the need to add elements to both the front and back of the sequence.

## What's Missing

You will also notice there are two member functions implemented in `vector` but not in `deque`, and, as you will see, `deque` doesn't need them.

1. `capacity()` - Returns the current capacity of a `vector`.
2. `reserve()` - Allocates room for a specified number of elements in a `vector`.

Herein lies the true beginning of our study. As it turns out, there is a stark difference between `vector` and `deque` in how they manage their internal storage under the hood. The `deque` allocates memory in chunks as it grows, with room for a fixed number of elements in each one. However, `vector` allocates its memory in contiguous blocks (which isn't necessarily a bad thing). But the interesting thing about `vector` is that the size of the internal buffer grows increasingly larger with each allocation after the `vector` realizes the current one isn't big enough. The following experiment sets out to prove why `deque` doesn't need `capacity()` or `reserve()` for that very reason.

## Experiment 1 - Container Growth

### Objective

The objective of this experiment is to observe the differences in container growth between the `vector` and the `deque`. The results of this experiment will illustrate these differences in terms of physical memory allocation and application performance.

### Description

The test application for this experiment is designed to read text from a file and use each line as the element to `push_back()` onto the `vector` and the `deque`. In order to generate large numbers of insertions, the file may be read more than once. The class to handle the test is shown below:

Hide Shrink ▲ Copy Code

```
#include <deque>
#include <fstream>
#include <string>
#include <vector>

static enum modes
{
    FM_INVALID = 0,
    FM_VECTOR,
    FM_DEQUE
};

class CVectorDequeTest
{
public:
    CVectorDequeTest();

    void ReadTestFile(const char* szFile, int iMode)
    {
        char buff[0xFFFF] = {0};
        std::ifstream inFile;
```

```
inFile.open(szFile);

while(!inFile.eof())
{
    inFile.getline(buff, sizeof(buff));

    if(iMode == FM_VECTOR)
        m_vData.push_back(buff);
    else if(iMode == FM_DEQUE)
        m_dData.push_back(buff);
}

inFile.close();

}

virtual ~CVectorDequeTest();

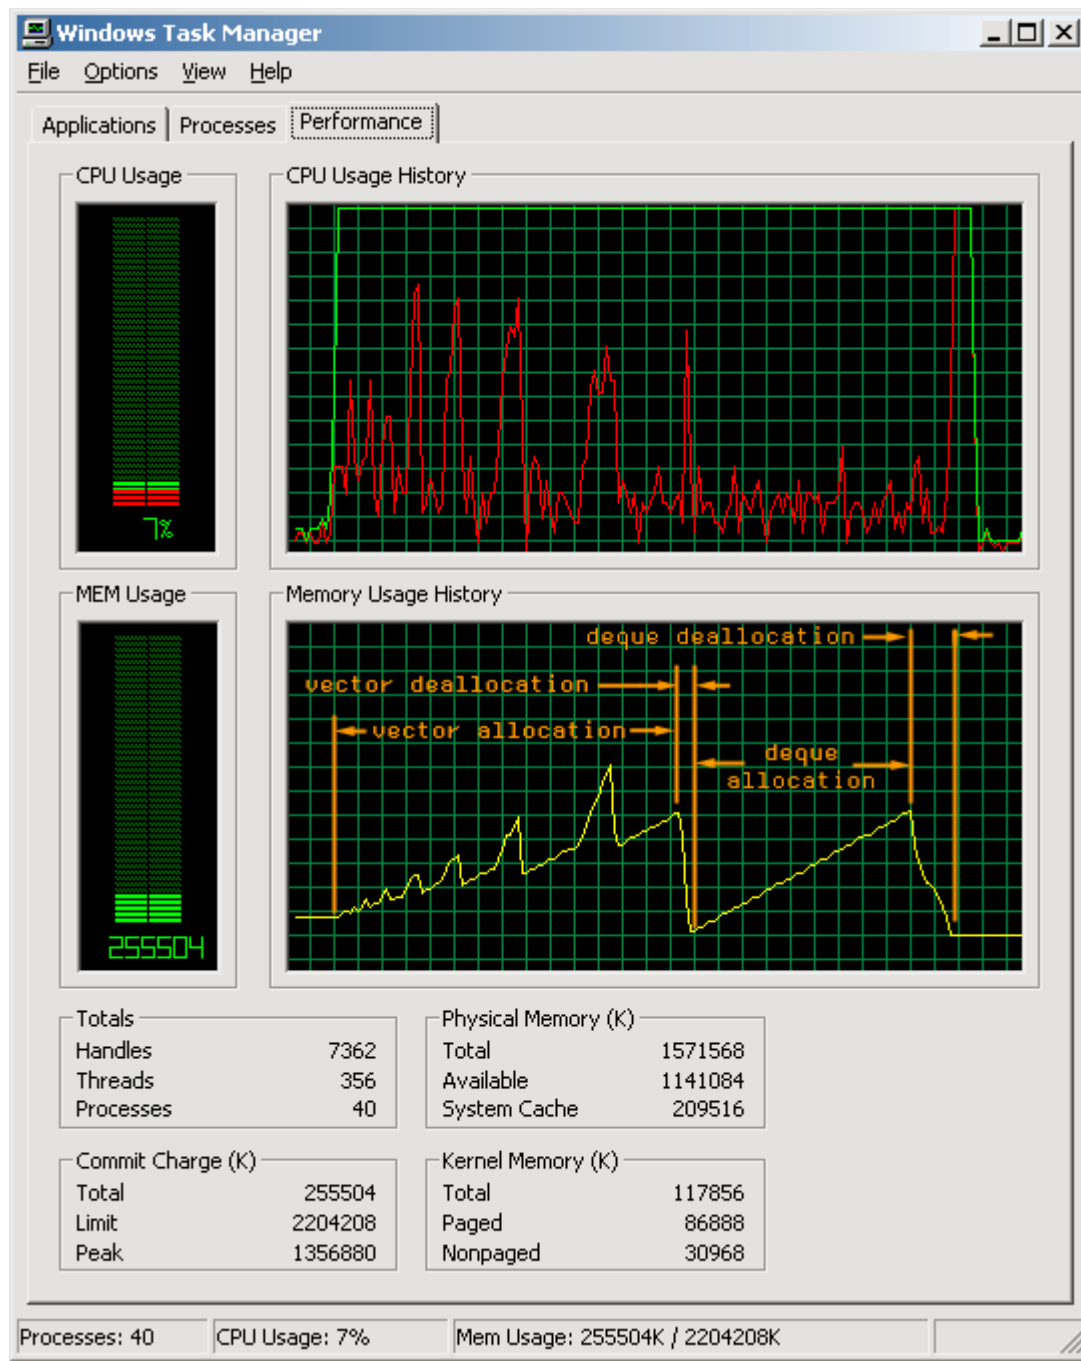
protected:
    std::vector<std::string> m_vData;
    std::deque<std::string> m_dData;
};
```

## Results

The test was performed under the following conditions:

Processor	1.8 GHz Pentium 4
Memory	1.50 GB
OS	W2K-SP4
No. of Lines in File	9874
Avg. Chars per Line	1755.85
No. of Times File Read	45
Total Elements Inserted	444330

The system performance was logged via Windows Task Manager, and the program was timed using Laurent Guinnard's [CDuration](#) class. The system performance graph is illustrated below:



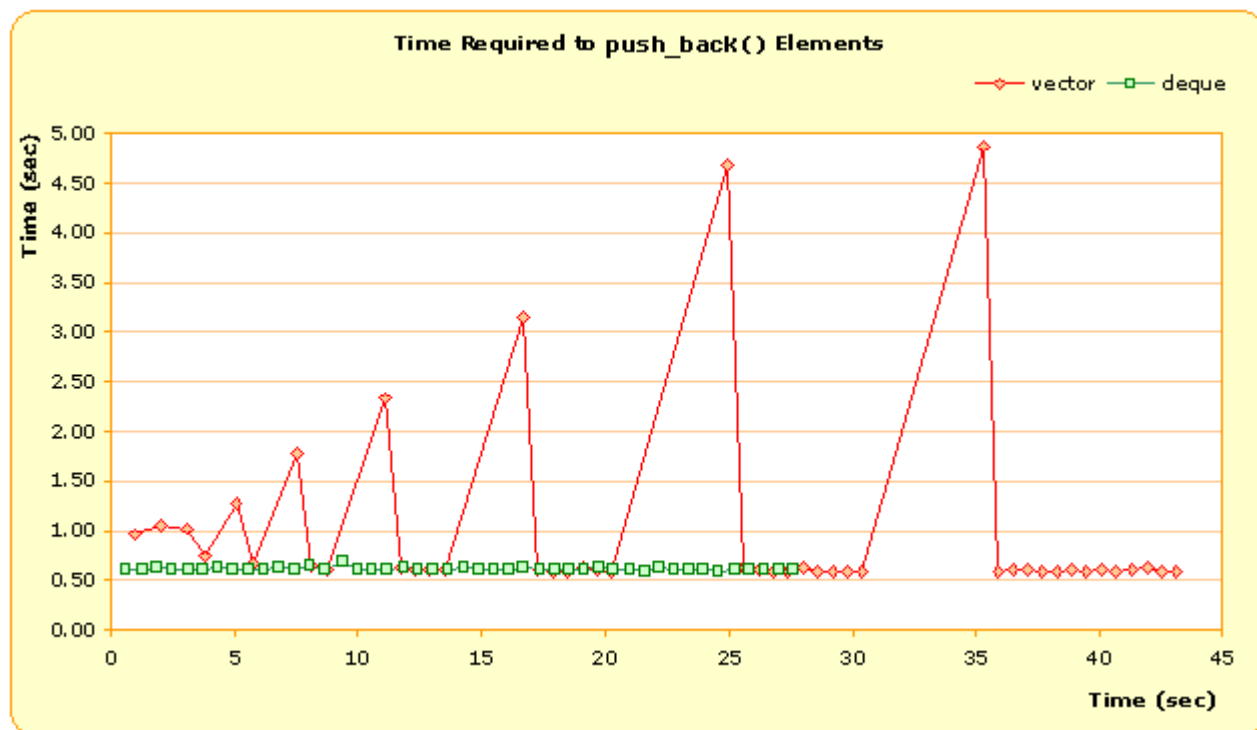
Note the peaks in memory usage during **vector** allocation, and how the peaks grow larger as **vector** allocates increasing internal buffer storage. Note also that **deque** does not exhibit this behavior, and the buffer continues to grow linearly with element insertion. The jump in kernel time during **deque** deallocation as well as the shape of the curve as memory is reclaimed was an unexpected result at first. I would have expected the deallocation to look similar to **vector**. After looking into things further and conducting some more tests, I was able to come up with a hypothesis: since **deque** memory is not contiguous, it must be more difficult to hunt down and reclaim. We will put this hypothesis to the test later, but first let's analyze the performance aspects of this experiment.

## Just how long do those memory allocations take?

Notice in the figure below that no elements were being added during the time **vector** was out finding more memory.



It is also of interest to notice how long each set of `push_back()` takes. This is illustrated in the figure below. Remember, each sample is 9874 strings added, with an average length of 1755.85.



## Experiment 2 - The Effects of `vector::reserve()`

### Objective

The objective of this experiment is to observe the benefits of calling `reserve()` on a `vector` before a large number of elements will be added and compare these results with `deque`, in terms of memory allocation and performance.

## Description

The test description for this experiment is the same as that of [Experiment 1](#), except that the following code was added to the test class constructor:

[Hide](#) [Copy Code](#)

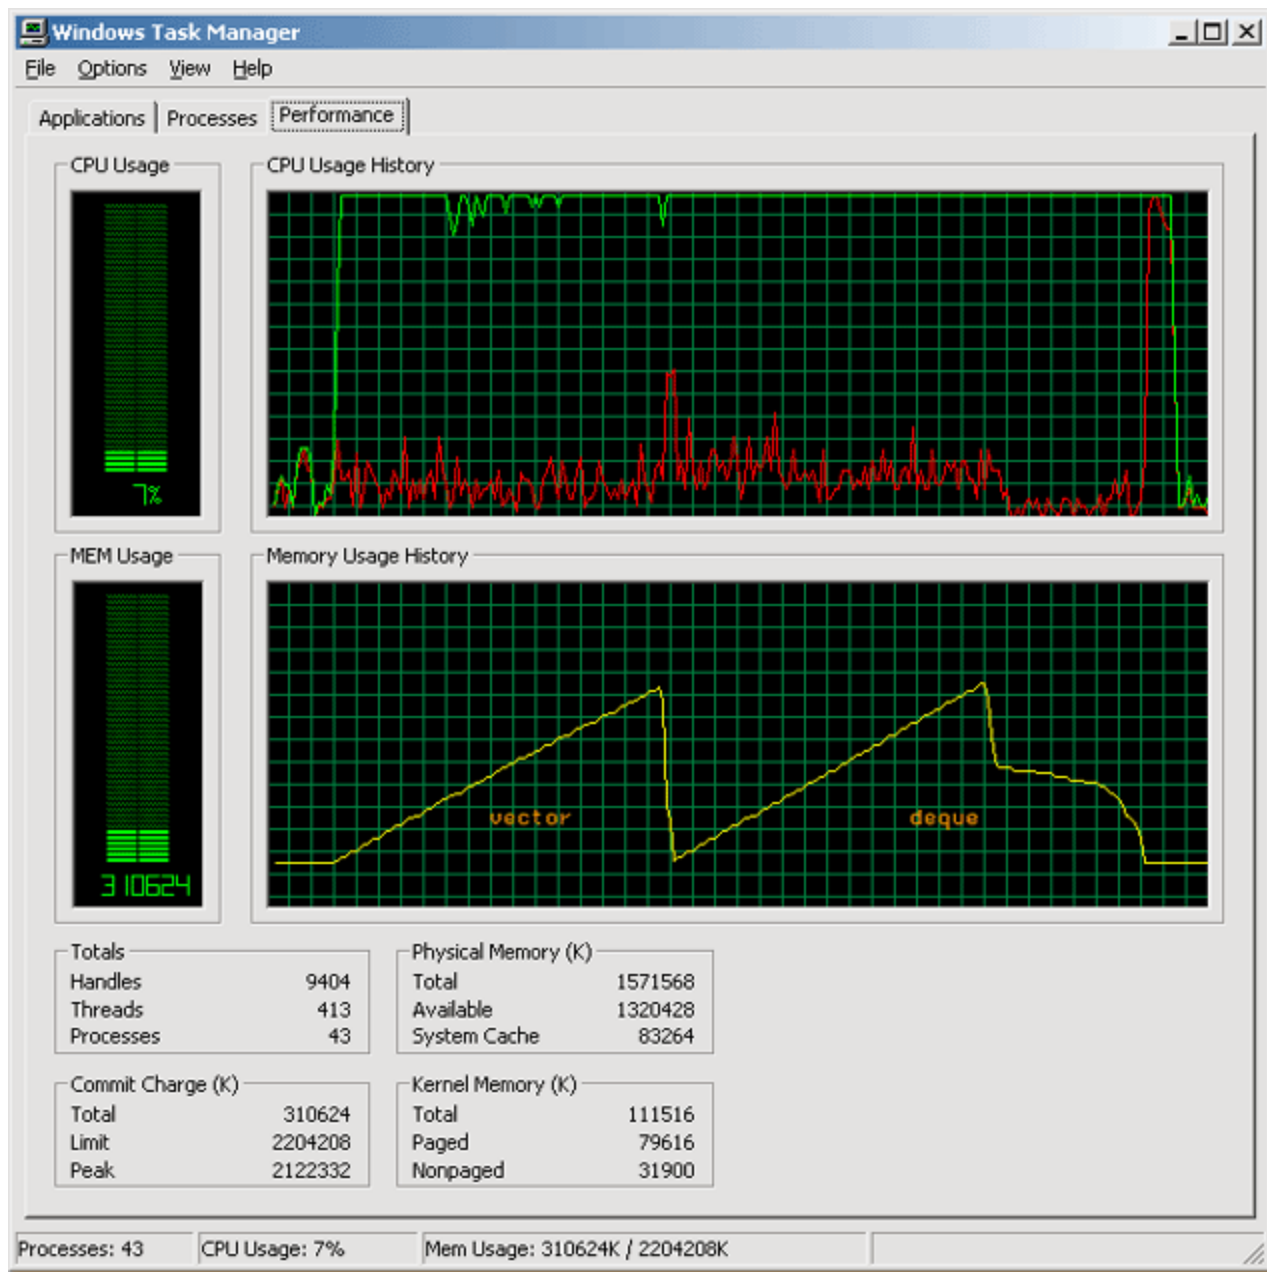
```
m_vData.reserve(1000000);
```

## Results

The test was performed under the following conditions:

Processor	1.8 GHz Pentium 4
Memory	1.50 GB
OS	W2K-SP4
No. of Lines in File	9874
Avg. Chars per Line	1755.85
No. of Times File Read	70
Total Elements Inserted	691180

The system performance was logged via Windows Task Manager, and the program was timed using Laurent Guinnard's [CDuration](#) class. The system performance graph is illustrated below:

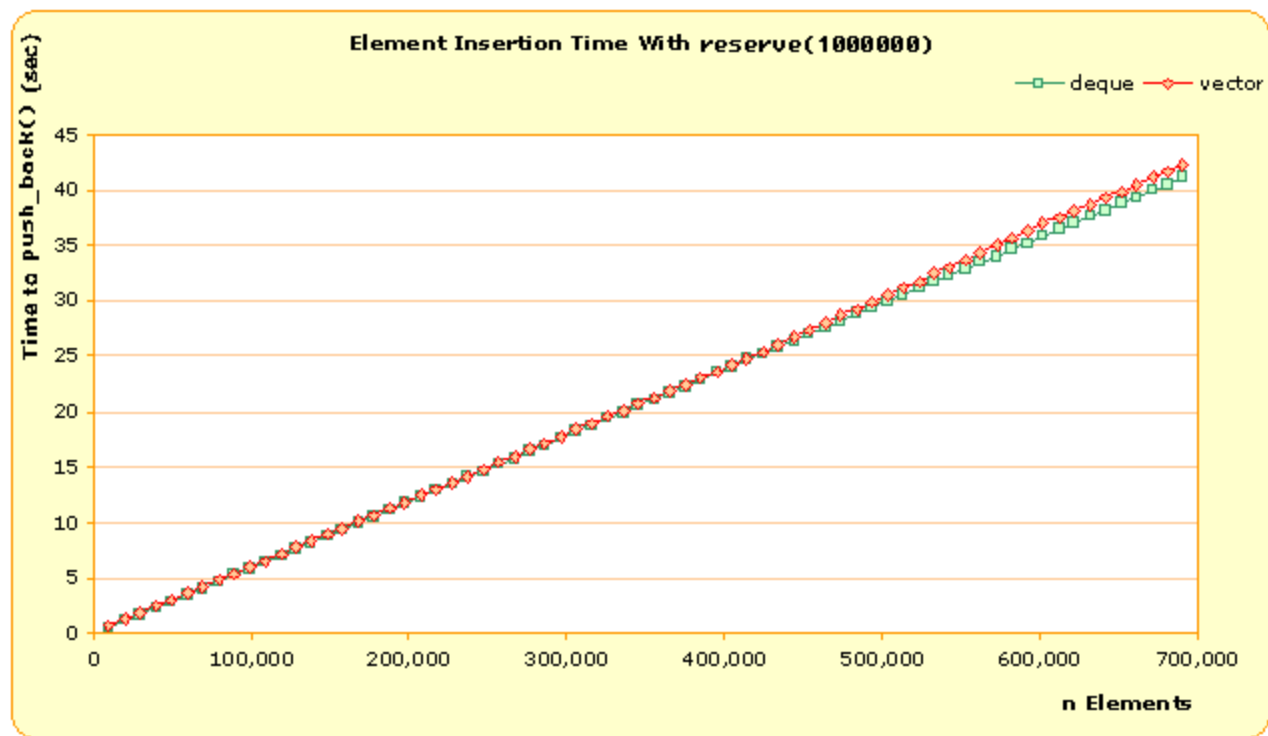


It is of interest to notice that **vector** no longer needs to allocate more internal buffer storage. The call to **reserve()** takes a single step to reserve more than enough space for our test platform of 691180 elements. As for the **deque** deallocation hypothesis, observe the drastic growth in memory deallocation time between this test and the previous one. We will quantify this in our next experiment.

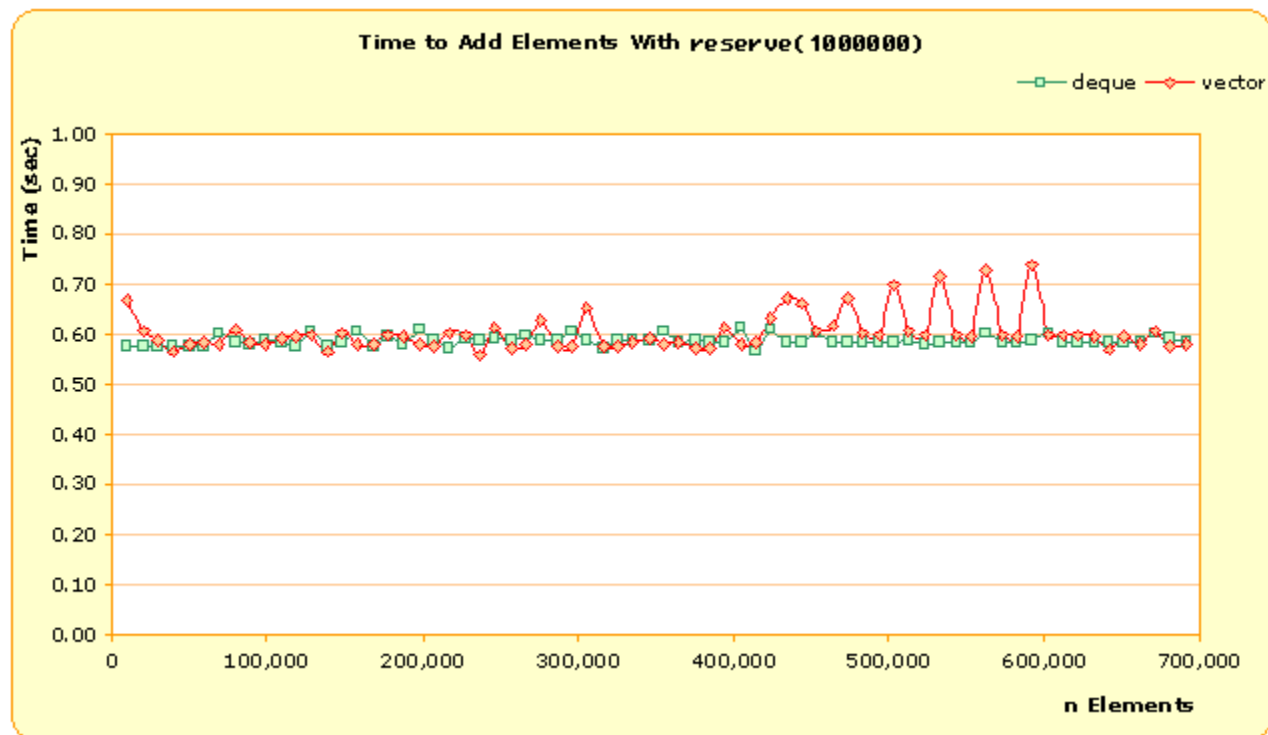
## How has this improved memory allocation performance?

The following figure illustrates the number of elements added to the containers over time:





As you can see, **vector** is now very close to **deque** in performance, when adding elements to the container. However, **vector** tends to be slightly more sporadic in how long it takes to insert a given set of elements. This is illustrated in the figure below:



A statistical analysis of the variability in **vector** vs. **deque**, with respect to the time it takes to insert 9874 elements of 1755.85 average length, is summarized in the following tables:

#### Vector

Mean 0.603724814 sec  
Maximum 0.738313000 sec  
Minimum 0.559959000 sec

#### Deque

Mean 0.588021114 sec  
Maximum 0.615617000 sec  
Minimum 0.567503000 sec

Std. Dev 0.037795736 sec  
6-Sigma 0.226774416 sec

Std. Dev 0.009907800 sec  
6-Sigma 0.059446800 sec

## Experiment 3 - Reclaiming Memory

### Objective

The objective of this experiment is to analyze and attempt to quantify the hypothesis that **deque** memory is more difficult to reclaim due to its non-contiguous nature.

### Description

The test class from [Experiment 1](#) will be utilized again in this experiment. The calling function is designed to allocate test classes of increasing size and log their performance accordingly. This implementation is as follows:

Hide Shrink ▲ Copy Code

```
for(xRun=0; xRun<NUMBER_OF_XRUNS; xRun++)
{
    df = new CVectorDequeTest;

    elapsed_time = 0;
    for(i=0; i<NUMBER_OF_RUNS*xRun; i++)
    {
        cout << "Deque - Run " << i << " of " << NUMBER_OF_RUNS*xRun << "... ";
        df->ReadTestFile("F:\\huge.csv",DF_DEQUE);

        deque_data.push_back(datapoint());

        deque_data.back().time_to_read = df->GetProcessTime();
        elapsed_time += deque_data.back().time_to_read;

        deque_data.back().elapsed_time = elapsed_time;

        cout << deque_data.back().time_to_read << " seconds\n";
    }

    vnElements.push_back(df->GetDequeSize());

    cout << "\n\nDeleting... ";

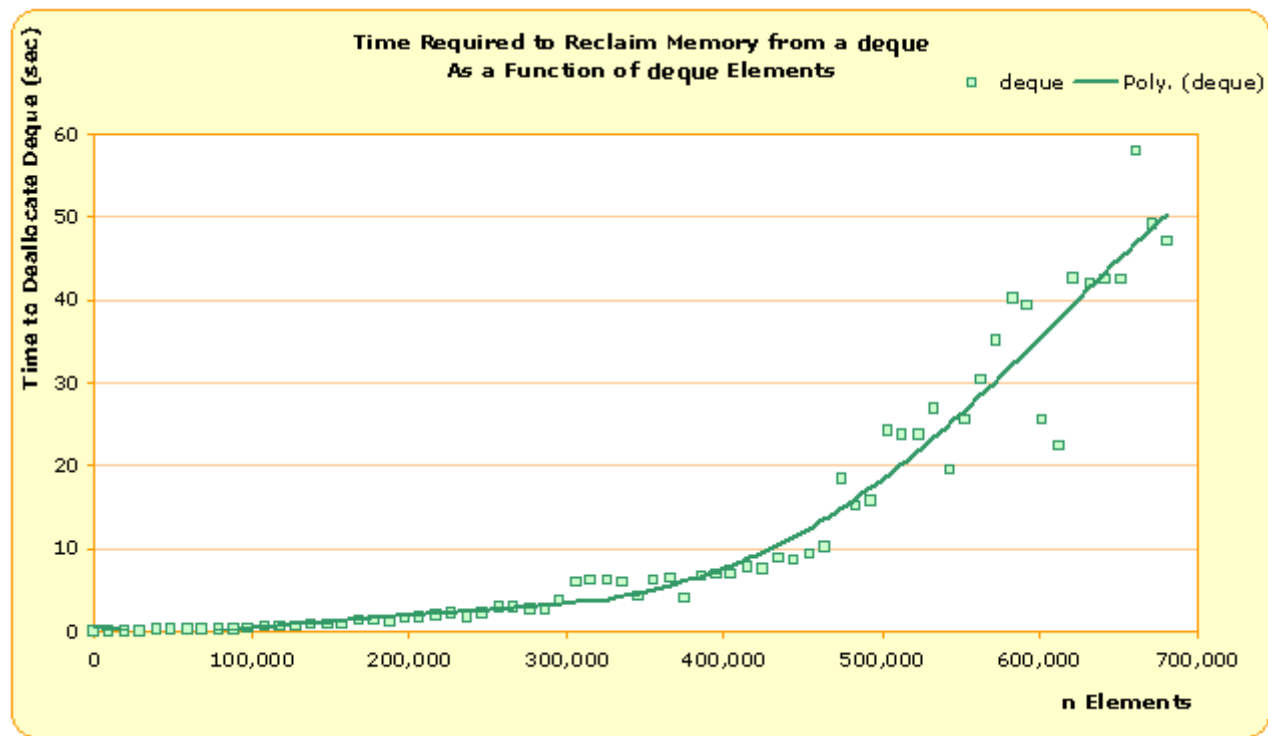
    del_deque.Start();
    delete df;
    del_deque.Stop();

    cout << del_deque.GetDuration()/1000000.0 << " seconds.\n\n";

    vTimeToDelete.push_back(del_deque.GetDuration()/1000000.0);
}
```

### Results

This experiment was performed on the same platform as the previous two experiments, except that the number of allocations was varied from 9874 to 691180 across 70 increments. The following figure illustrates the time required to reclaim **deque** memory as a function of the number of elements in the **deque**. The **deque** was filled with strings with an average length of 1755.85 chars.



Although the actual time varies significantly from the trendline in several instances, the trendline holds accurate with an  $R^2=95.15\%$ . The actual deviation of any given data point from the trendline is summarized in the following table:

#### deque Results

Mean 0.007089269 sec  
Maximum 11.02838496 sec  
Minimum -15.25901667 sec  
Std. Dev 3.3803636 sec  
6-Sigma 20.2821816 sec

This is fairly significant when compared to the results of **vector** in the same scenario. The following figure shows deallocation times for **vector** under the same loading as **deque** above:

The data in this test holds an  $R^2=81.12\%$ . This could likely be improved with more iterations of each data point and averaging the runs. Nonetheless, the data is suitable to mark the point in question, and the deviation of any given data point from the trendline is summarized in the following statistical parameters:

#### vector Results

Mean     -0.007122715 sec  
Maximum 0.283452127 sec  
Minimum -0.26724459 sec  
Std. Dev 0.144572356 sec  
6-Sigma 0.867434136 sec

## Experiment 4 - Performance Characteristics of vector::insert() vs. deque::insert()

### Objective

The "claim to fame" as it were for **deque** is the promise of constant-time **insert()**. Just how does this stack up against **vector::insert()**? The objective of this experiment is (not surprisingly) to observe the performance characteristics of **vector::insert()** vs. **deque::insert()**.

### Description

There may be times when adding things to the back of a container doesn't quite suit your needs. In this case, you may want to employ **insert()**. This experiment also has the same form as [Experiment 1](#), however instead of doing **push\_back()**, the test does **insert()**.

### Results

As you can see in the following figures, the benefit of constant-time insertion offered by **deque** is staggering when compared against **vector**.

Note the difference in time-scales, as 61810 elements were added to these containers.

## Experiment 5 - Container Retrieval Performance

### Objective

This experiment will test the performance of `vector::at()`, `vector::operator[]`, `deque::at()` and `deque::operator[]`. It has been suggested that `operator[]` is faster than `at()` because there is no bounds checking, also it has been requested to compare `vector` vs. `deque` in this same regard.

## Description

This test will insert 1000000 elements of type `std::string` with a length of 1024 characters into each container and measure how long it takes to access them all via `at()` and `operator[]`. The test will be performed 50 times for each scenario and the results presented as a statistical summary.

## Results

Well, perhaps surprisingly, there is very little difference in performance between `vector` and `deque` in terms of accessing the elements contained in them. There is also negligible difference between `operator[]` and `at()` as well. These results are summarized below:

<code>vector::at()</code>	<code>deque::at()</code>
Mean 1.177088125 sec	Mean 1.182364375 sec
Maximum 1.189580000 sec	Maximum 1.226860000 sec
Minimum 1.168340000 sec	Minimum 1.161270000 sec
Std. Dev 0.006495193 sec	Std. Dev 0.016362148 sec
6-Sigma 0.038971158 sec	6-Sigma 0.098172888 sec
<code>vector::operator[]</code>	<code>deque::operator[]</code>
Mean 1.164221042 sec	Mean 1.181507292 sec
Maximum 1.192550000 sec	Maximum 1.218540000 sec
Minimum 1.155690000 sec	Minimum 1.162710000 sec
Std. Dev 0.007698520 sec	Std. Dev 0.010275712 sec
6-Sigma 0.046191120 sec	6-Sigma 0.061654272 sec

## Conclusions

In this article, we have covered several different situations where one could possibly have a need to choose between `vector` and `deque`. Let's summarize our results and see if our conclusions are in line with the [standard](#).

When performing a large number of `push_back()` calls, remember to call `vector::reserve()`.

In [Experiment 1](#), we studied the behavior of container growth between `vector` and `deque`. In this scenario, we saw that since `deque` allocates its internal storage in blocks of pre-defined size, `deque` can grow at a constant rate. The performance of `vector` in this experiment then led us to think about calling `vector::reserve()`. This was then the premise for [Experiment 2](#), where we basically performed the same experiment except that we had called `reserve()` on our `vector`. This then is grounds for holding on to `vector` as our default choice.

If you are performing many deallocations, remember that `deque` takes longer to reclaim memory than `vector`.

In [Experiment 3](#), we explored the differences between reclaiming the contiguous and non-contiguous memory blocks of `vector` and `deque`, respectively. The results proved that `vector` reclaims memory in linear proportion to the number of elements whereas

`deque` is exponential. Also, `vector` is several orders of magnitude than `deque` in reclaiming memory. As a side note, if you are performing your calls to `push_back()` within a tight loop or sequence, there is a significant possibility that most of the memory `deque` obtains, will be contiguous. I have tested this situation for fun and have found the deallocation time to be close to `vector` in these cases.

If you are planning to use `insert()`, or have a need for `pop_front()`, use `deque`.

Well, ok, `vector` doesn't have `pop_front()`, but based on the results of [Experiment 4](#), it might as well not have `insert()` either. The results of [Experiment 4](#) speak volumes about the need for `deque` and why it is part of the STL as a separate container class.

For element access, `vector::at()` wins by a nose.

After summing up the statistics of [Experiment 5](#), I would have to say that although all the methods were close, `vector::at()` is the winner. This is because of the best balance between the raw mean of the access times as well as the lowest 6-sigma value.

What's all this 6-Sigma stuff?

Although a popular buzzword in industry today, 6-Sigma actually has its roots in statistics. If you generate a Gaussian distribution (or Bell-curve) for your sampled data, you can show that at one standard deviation (the symbol for std. deviation is the Greek letter, sigma, BTW) from the mean, you will have 68.27% of the area under the curve covered. At 2 Standard Deviations, 2-Sigma, you have 95.45% of the area under the curve, at 3 Standard Deviations, you will have 99.73% of the area and so forth until you get to 6 standard deviations, when you have 99.99985% of the area (1.5 Defects per million, 6-Sigma).

## Final Words

I hope you have gained some insight into `deque` and have found this article both interesting and enlightening. Any questions or comments are certainly welcome and any discussion on `vector` or `deque` is encouraged.

## References

1. Plauger, P.J. [Standard C++ Library Reference](#). February, 2003. MSDN.
2. ISO/IEC 14882:1998(E). [Programming Languages - C++](#). ISO and ANSI C++ Standard.
3. Schildt, Herbert. [C++ from the Ground Up, Second Edition](#). Berkeley: 1998.
4. Sutter, Herb. [More Exceptional C++](#). Indianapolis: 2002.

## License

This article has no explicit license attached to it but may contain usage terms in the article text or the download files themselves. If in doubt please contact the author via the discussion board below.

A list of licenses authors might use can be found [here](#)

## Share

## About the Author



### Nitron



Software Developer (Senior) Lockheed Martin  
United States

Walter Storm is currently doing quantitative research and data science. Originally from Tunkhannock, PA., he has a B.S. in Aerospace Engineering from [Embry-Riddle Aeronautical University](#)<sup>[^]</sup>, and an M.S. in Systems Engineering from [SMU](#)<sup>[^]</sup>. He has been professionally developing software in some form or another since January of 2001.

[Permalink](#)

Layout: [fixed](#) | [fluid](#)

[View Walter Storm's profile on LinkedIn.](#)<sup>[^]</sup>

[Privacy](#)

[Cookies](#)

[Terms of Use](#)

Article Copyright 2003 by Nitron  
Everything else Copyright © [CodeProject](#),  
1999-2020

Web05 2.8.20201103.1

## Comments and Discussions

You must [Sign In](#) to use this message board.

Search Comments

Spacing

Relaxed



Layout

Normal



Per page

25


































































[Update](#)

First Prev Next

<a href="#">article</a>	Member 13195584	11-May-17 21:46
<a href="#">Experiment 3 - Exponential time to deallocate memory for a deque</a>	milo234	18-May-16 19:04
<a href="#">Re: Experiment 3 - Exponential time to deallocate memory for a deque</a>	Nitron	25-Sep-17 4:49
<a href="#">erase performance</a>	GHop	7-Feb-14 10:24



 <b>Experiment 5 strange results (Container Retrieval Performance)</b> 	 <b>Jean-Marc Dressler</b>	<b>9-Dec-10 0:49</b>
 <b>I do not agree with your analysis on experiment 3</b> 	 <b>Shinhy Ku</b>	<b>25-Oct-08 6:15</b>
 Re: I do not agree with your analysis on experiment 3 	 Nitron	31-May-10 17:57
 <b>Does it mean that we should prefer deque over vector</b> 	 <b>yccheok</b>	<b>14-Sep-06 1:22</b>
 Re: Does it mean that we should prefer deque over vector 	 Nitron	27-Sep-06 18:23
 <b>Very nice and...</b> 	 <b>Muzero2</b>	<b>7-Feb-06 0:48</b>
 <b>6 Sigma explanation Error</b> 	 <b>Steve Chernyavskiy</b>	<b>8-Jun-05 8:08</b>
 Re: 6 Sigma explanation Error 	 Nitron	13-Jun-05 13:07
 <b>Very Usefull</b> 	 <b>Sudhir Mangla</b>	<b>9-Mar-05 18:40</b>
 Re: Very Usefull 	 Nitron	13-Jun-05 13:07
 <b>Well done</b> 	 <b>nohoper</b>	<b>26-Dec-04 20:53</b>
 <b>Great article</b> 	 <b>Simon Hughes</b>	<b>5-Oct-04 12:01</b>
 Re: Great article 	 Nitron	13-Jun-05 13:09
 <b>Thank you</b> 	 <b>mikeboedi</b>	<b>28-Mar-04 21:54</b>
 <b>But deque has linear insertion time?</b> 	 <b>Eloff</b>	<b>29-Jan-04 7:13</b>
 Re: But deque has linear insertion time? 	 Nitron	2-Feb-04 11:20
 Re: But deque has linear insertion time? 	 Eloff	3-Feb-04 9:31
 Re: But deque has linear insertion time? 	 EnderJSC	15-Mar-04 12:14
 Re: But deque has linear insertion time? 	 Eloff	19-Mar-04 9:12
 Re: But deque has linear insertion time? 	 <b>Anonymous</b>	26-May-04 7:35
 Re: But deque has linear insertion time? 	 Albrecht Fritzsche	3-Dec-04 1:34

Last Visit: 3-Nov-20 19:08    Last Update: 3-Nov-20 19:08

[Refresh](#)[1](#) [2](#) [3](#) [Next >](#)  
General  
News



Suggestion



Question



Bug



Answer



Joke



Praise



Rant



Admin

Use  
Ctrl+Left/Right  
to  
switch  
messages,  
Ctrl+Up/Down  
to  
switch  
threads,  
Ctrl+Shift+Left/Right  
to  
switch  
pages.